

# CSCI E-89C Deep Reinforcement Learning

Harvard Summer School

Dmitry Kurochkin

Summer 2020

Lecture 5

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_\pi(s)$
    - Off-policy Estimation of  $q_\pi(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_*(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Target Policy v.s. Behavior Policy

- 1 The goal is to learn *target* policy  $\pi(a|s)$ , usually deterministic greedy policy
- 2 Off-policy methods estimate *target* policy using data generated according to *behavior* policy  $b(a|s)$
- 3 Off-policy methods usually have larger variance and slower convergence
- 4 On-policy learning, i.e. whenever  $b(a|s) = \pi(a|s)$ , is a special case of off-policy
- 5 Off-policy methods can be used to learn from available data generated by a non-learning controller
- 6 Assumption of coverage:  $\pi(a|s) > 0 \Rightarrow b(a|s) > 0$

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_\pi(s)$
    - Off-policy Estimation of  $q_\pi(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_*(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Importance Sampling

If we start in state  $S_t$  and follow policy  $\pi$ , then

$$\begin{aligned}
 &P_{\pi}\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t\} \\
 &= P_{\pi}\{A_t | S_t\} P_{\pi}\{S_{t+1}, A_{t+1}, \dots, S_T | S_t, A_t\} \\
 &= P_{\pi}\{A_t | S_t\} P_{\pi}\{S_{t+1} | S_t, A_t\} P_{\pi}\{A_{t+1}, \dots, S_T | S_t, A_t, S_{t+1}\} \\
 &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) P_{\pi}\{A_{t+1}, \dots, S_T | S_{t+1}\} \\
 &\vdots \\
 &= \pi(A_t | S_t) p(S_{t+1} | S_t, A_t) \pi(A_{t+1} | S_{t+1}) \dots p(S_T | S_{T-1}, A_{T-1}) \\
 &= \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)
 \end{aligned}$$

# Importance Sampling

For any policy  $\pi(a|s)$ :

$$P_{\pi}\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t\} = \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)$$

# Importance Sampling

For any policy  $\pi(a|s)$ :

$$P_{\pi}\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t\} = \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)$$

Define the importance-sampling ratio:

$$\begin{aligned} \rho_{t:(T-1)} &\doteq \frac{P_{\pi}\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t\}}{P_b\{A_t, S_{t+1}, A_{t+1}, \dots, S_T | S_t\}} \\ &= \frac{\prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1} | S_k, A_k)}{\prod_{k=t}^{T-1} b(A_k | S_k) p(S_{k+1} | S_k, A_k)} \\ &= \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k)}{b(A_k | S_k)} \end{aligned}$$

Given trajectories,  $\rho_{t:(T-1)}$  depends on the policies only!



# Importance Sampling

We notice that under policy  $b$ ,

$$E_b [G_t | S_t = s] = v_b(s),$$

but the expected transformed cumulative discounted return is

$$E_b [\rho_{t:(T-1)} G_t | S_t = s] = E_\pi [G_t | S_t = s] = v_\pi(s),$$

where data were generated under  $b$ .

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Off-policy Estimation of $v_\pi(s)$

Let index  $t$  run through episodes. In order to estimate  $v_\pi(s)$ , the agent can follow policy  $b \neq \pi$  but transform the observations for  $G_t$ 's as follows:

- 1 Ordinary importance sampling (unbiased in case of first-visit MC):

$$V(s) \doteq \frac{\sum_{t \in \mathcal{T}(s)} \rho_{t:(T-1)} G_t}{|\mathcal{T}(s)|}.$$

- 2 Weighted importance sampling (biased):

$$V(s) \doteq \begin{cases} \frac{\sum_{t \in \mathcal{T}(s)} \rho_{t:(T-1)} G_t}{\sum_{t \in \mathcal{T}(s)} \rho_{t:(T-1)}}, & \text{if } \sum_{t \in \mathcal{T}(s)} \rho_{t:(T-1)} \neq 0, \\ 0, & \text{otherwise.} \end{cases}$$

Here,  $\mathcal{T}(s)$  is either

- (a) set of all time steps  $t$  in which state  $s$  is first visited (first-visit MC); or
- (b) set of all time steps  $t$  in which state  $s$  is visited (every-visit MC).

# Off-policy Estimation of $q_\pi(s, a)$

Let index  $t$  run through episodes. In order to estimate  $q_\pi(s, a)$ , the agent can follow policy  $b \neq \pi$  but transform the observations for  $G_t$ 's as follows:

- 1 Ordinary importance sampling (unbiased in case of first-visit MC):

$$Q(s, a) \doteq \frac{\sum_{t \in \mathcal{T}(s, a)} \rho_{t:(T-1)} G_t}{|\mathcal{T}(s, a)|}.$$

- 2 Weighted importance sampling (biased):

$$Q(s, a) \doteq \begin{cases} \frac{\sum_{t \in \mathcal{T}(s, a)} \rho_{t:(T-1)} G_t}{\sum_{t \in \mathcal{T}(s, a)} \rho_{t:(T-1)}}, & \text{if } \sum_{t \in \mathcal{T}(s, a)} \rho_{t:(T-1)} \neq 0, \\ 0 & , \text{ otherwise.} \end{cases}$$

Here,  $\mathcal{T}(s, a)$  is either

- (a) set of all time steps  $t$  in which the pair  $(s, a)$  is first visited (first-visit MC); or
- (b) set of all time steps  $t$  in which the pair  $(s, a)$  is visited (every-visit MC).

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Incremental Implementation

Let's fix pair  $(s, a)$ . We need to estimate  $q(s, a)$  from the sequence of returns

$$G_1, G_2, \dots, G_k, \dots, G_{n-1},$$

where  $k$  represents the  $k$ -th visit to pair  $(s, a)$ :

$$Q_n \doteq \frac{\sum_{k=1}^{n-1} w_k G_k}{\sum_{k=1}^{n-1} w_k}, n \geq 2.$$

Here,  $w_k = \rho_{t_k:(T(t_k)-1)}$  are the corresponding weights, where  $t_k$  is the time step in which  $(s, a)$  is visited  $k$ -th time.

The the updating rule is

$$Q_{n+1} \doteq Q_n + \frac{w_n}{C_n} [G_n - Q_n], n \geq 1$$

$$C_{n+1} \doteq C_n + w_{n+1}$$

# Off-policy Estimation of $\pi_*(s, a)$

## Off-policy MC control, for estimating $\pi \approx \pi_*$

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \in \mathbb{R}$  (arbitrarily)

$C(s, a) \leftarrow 0$

$\pi(s) \leftarrow \operatorname{argmax}_a Q(s, a)$  (with ties broken consistently)

Loop forever (for each episode):

$b \leftarrow$  any soft policy

Generate an episode using  $b$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

$W \leftarrow 1$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$  (with ties broken consistently)

If  $A_t \neq \pi(S_t)$  then exit inner Loop (proceed to next episode)

$W \leftarrow W \frac{1}{b(A_t|S_t)}$

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA



## One-step TD: Estimating $v_\pi(s)$

Recall that the state-value function is defined as follows:

$$v_\pi(s) \doteq E_\pi [G_t | S_t = s].$$

Every-visit MC method for nonstationary environment (*constant*- $\alpha$  MC):

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)], \quad \alpha \in (0, 1],$$

i.e. need to wait until the end of the episode because of  $G_t$ .

## One-step TD: Estimating $v_\pi(s)$

Recall that the state-value function is defined as follows:

$$v_\pi(s) \doteq E_\pi [G_t | S_t = s].$$

Every-visit MC method for nonstationary environment (*constant- $\alpha$*  MC):

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)], \quad \alpha \in (0, 1],$$

i.e. need to wait until the end of the episode because of  $G_t$ .

Also, recall that

$$v_\pi(s) = E_\pi [R_{t+1} + \gamma v_\pi(S_{t+1}) | S_t = s].$$

Then at time  $t + 1$  can update as follows (*one-step* TD or TD(0)):

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)], \quad \alpha \in (0, 1],$$

i.e. no need to wait until the end of the episode!

# One-step TD: Estimating $v_\pi(s)$

*One-step* TD prediction:

$$V(S_t) \leftarrow V(S_t) + \alpha \underbrace{[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]}_{\doteq \delta_t}.$$

TD error,  $\delta_t$ , is available at time time  $t + 1$ .

# One-step TD: Estimating $v_\pi(s)$

One-step TD prediction:

$$V(S_t) \leftarrow V(S_t) + \alpha \underbrace{[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]}_{\doteq \delta_t}.$$

TD error,  $\delta_t$ , is available at time time  $t + 1$ .

Note that the MC estimate  $V(S_t)$  does not change over the episode and the MC error is related to  $\delta_t$  as follows:

$$\begin{aligned} G_t - V(S_t) &= R_{t+1} + \gamma G_{t+1} - V(S_t) + \gamma V(S_{t+1}) - \gamma V(S_{t+1}) \\ &= [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] + \gamma (G_{t+1} - V(S_{t+1})) \\ &= \delta_t + \gamma (G_{t+1} - V(S_{t+1})) \\ &\vdots \\ &= \delta_t + \gamma \delta_{t+1} + \dots + \gamma^{T-t-1} \delta_{T-1} = \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k. \end{aligned}$$

# One-step TD: Estimating $v_{\pi}(s)$

*One-step* TD prediction:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] .$$

Algorithm:

## Tabular TD(0) for estimating $v_{\pi}$

Input: the policy  $\pi$  to be evaluated

Algorithm parameter: step size  $\alpha \in (0, 1]$

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

    until  $S$  is terminal

# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

# Advantages of TD Methods

*One-step* TD prediction:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)] .$$

- 1 Model-free
- 2 No need to wait until the end of the episode - TD may converge substantially faster
- 3 Can be applied to continuing tasks with no episodes
- 4 MC must discount some of the episodes - TD may converge substantially faster
- 5 Given the step-size parameter  $\alpha$  is sufficiently small, TD(0) converges to  $v_\pi$

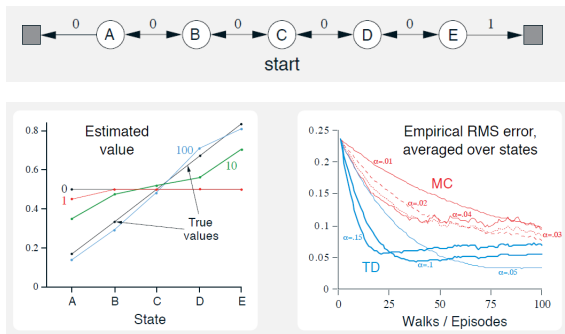
# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_{\pi}(s)$
    - Off-policy Estimation of  $q_{\pi}(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_{*}(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA



# Random Walk Example

In the example below the only action in all states is "wait." All transitions are equally likely. When the random walk terminates on the right, the reward is 1, otherwise all rewards are 0.



Left: TD(0) with  $\alpha = 0.1$  results for 0, 1, 10, and 100 episodes.

Right: Learning curves for MC and TD(0) and various step-size parameters  $\alpha$ .


# Contents

- 1 Off-policy Learning
  - Target Policy v.s. Behavior Policy
  - Importance Sampling
  - Off-policy Prediction
    - Off-policy Estimation of  $v_\pi(s)$
    - Off-policy Estimation of  $q_\pi(s, a)$
  - Off-policy Control
    - Incremental Implementation
    - Off-policy Estimation of  $\pi_*(s, a)$
- 2 One-step Temporal-Difference (TD) Learning
  - TD Prediction
  - Advantages of TD Methods
  - Example
  - TD Control: SARSA

## SARSA: Estimation of $q_*(s, a)$

Similarly to TD prediction, the updating rule for  $Q(S_{t+1}, A_{t+1})$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$



# SARSA: Estimation of $q_*(s, a)$

Similarly to TD prediction, the updating rule for  $Q(S_{t+1}, A_{t+1})$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

    until  $S$  is terminal