

# CSCI S-89C Deep Reinforcement Learning

Harvard Summer School

Dmitry Kurochkin

Summer 2020

Lecture 9

# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Quiz 8

## Question 1

4 / 4 pts

TD( $\lambda$ ) with  $\lambda = 0$  is equivalent to

- (A) 1-step TD
- (B) Monte Carlo Control
- (C) Monte Carlo Prediction
- (D) None of (A), (B), (C)

Please select:

☐ B☐ C☒ A☐ D

Correct!

# Quiz 8

## Question 2

4 / 4 pts

SARSA( $\lambda$ ) with  $\lambda = 0$  is equivalent to

- (A) SARSA
- (B) Monte Carlo Control
- (C) Monte Carlo Prediction
- (D) None of (A), (B), (C)

Please select:

Correct!

☒ A

☐ C

☐ B

☐ B

# Quiz 8

## Question 3

4 / 4 pts

Q-learning is equivalent to SARSA( $\lambda$ ) for some choice of parameter  $\lambda \in [0, 1]$ .

☐ True☒ False**Correct!**

# Quiz 8

## Question 4

4 / 4 pts

The environment has three states:  $s_A$ ,  $s_B$ , and  $s_C$ . In each state there are two actions,  $a_1$  and  $a_2$ , available.

Suppose we want to obtain optimal policy using Q-learning; and the initial values of the action-value function are

$$Q(s_A, a_1) = 6, Q(s_A, a_2) = 5,$$

$$Q(s_B, a_1) = 4, Q(s_B, a_2) = 3,$$

$$Q(s_C, a_1) = 2, Q(s_C, a_2) = 1.$$

We generate the sequence using  $\epsilon$ -greedy ( $\epsilon = 0.05$ ) policy with respect to current  $Q(s, a)$  values and observe:

$$s_A, a_1, R_1 = 5, s_B, a_2, \dots$$

If  $\alpha = 0.1, \gamma = 0.9$ , what is  $Q(s_A, a_1)$  after its first update according to the Q-learning algorithm?

Hint: use the following off-policy Q-learning:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)].$$

Correct!

6.26

Correct Answers

6.26 (with margin: 0.01)

# Quiz 8

## Question 5

4 / 4 pts

Suppose that the environment has three states {stage  $k$ } with  $k=1, 2, 3$  ( $k$  is a *feature* of the state):

$$s_A = \{\text{stage } 1\}, s_B = \{\text{stage } 2\}, s_C = \{\text{stage } 3\}.$$

In each state there are two actions, 0 and 1, available.

The action-value function is being approximated by the following linear function (linear in weights):

$$q_\pi(s, a) \approx \hat{q}(s, a, \mathbf{w}) \doteq w_1 + w_2 a + (w_3 + w_4 a) k,$$

where  $\mathbf{w} = (w_1, w_2, w_3, w_4)^T$  are weights and  $k$  corresponds to the state  $s = \{\text{stage } k\}$ .

$$\text{For example, } q_\pi(s_C, 0) \approx \hat{q}(s_C, 0, \mathbf{w}) \doteq w_1 + w_2 \cdot 0 + (w_3 + w_4 \cdot 0) \cdot 3.$$

Here,  $\pi$  is some current policy, usually  $\epsilon$ -greedy with respect to the most recent  $\hat{q}$ .

Suppose we run Q-learning with approximation:

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha [R_{t+1} + \gamma \max_a \hat{q}(S_{t+1}, a, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t)] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t),$$

where the initial values (i.e. values at time  $t = 0$ ) for weights are set to zero,  $\mathbf{w}_0 = (0, 0, 0, 0)^T$ , and observe:

$$s_B, A_0 = 1, R_1 = 5, s_C, A_1 = 0, \dots$$

If  $\alpha = 0.1, \gamma = 0.9$ , what is  $\mathbf{w}_1$  after the first update (i.e. weights at time  $t = 1$ )?

Please select:

☐ (0.5, 0.5, 1.0, 0)

☒ (0.5, 0.5, 1.0, 1.0)

☐ (0.5, 0, 1.0, 0)

☐ (5, 0, 10, 0)

Correct!



# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Mini-batch Gradient Descent

Given policy  $\pi$ , assume that for some weights  $\mathbf{w} = (w_1, w_2, \dots, w_d)^T$  (usually  $d \ll |\mathcal{S}|$ ) we can approximate:

$$v_{\pi}(s) \approx \hat{v}(s, \mathbf{w}).$$

# Mini-batch Gradient Descent

Given policy  $\pi$ , assume that for some weights  $\mathbf{w} = (w_1, w_2, \dots, w_d)^T$  (usually  $d \ll |\mathcal{S}|$ ) we can approximate:

$$v_\pi(s) \approx \hat{v}(s, \mathbf{w}).$$

The Mini-batch gradient descent (SGD) method that minimizes the mean-squared error

$$J(\mathbf{w}) \doteq E_\pi \left[ (v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}))^2 \right]$$

is then

$$\begin{aligned} \mathbf{w}_{k+1} &\doteq \mathbf{w}_k - \frac{1}{2} \alpha \nabla \sum_{t=mk}^{m(k+1)-1} [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_k)]^2 \\ &= \mathbf{w}_k + \alpha \sum_{t=mk}^{m(k+1)-1} [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_k)] \nabla \hat{v}(S_t, \mathbf{w}_k), \end{aligned}$$

where  $m$  is the size of *mini-batches*.


# Mini-batch Gradient Descent

Since we do not know  $v_\pi(S_t)$ , we use an approximation  $U_t$  of the state value function (for example  $G_t$  in case of MC). The weights then can be obtained as follows:

$$\mathbf{w}_{k+1} \doteq \mathbf{w}_k + \alpha \sum_{t=mk}^{m(k+1)-1} \left[ \underbrace{U_t}_{\approx v_\pi(S_t)} - \hat{v}(S_t, \mathbf{w}_k) \right] \nabla \hat{v}(S_t, \mathbf{w}_k)$$

# Mini-batch Gradient Descent

Example: Path in  $(w_1, w_2)$  plane:



Source: *Hands-On Machine Learning with Scikit-Learn and TensorFlow* by A. Géron

# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Supervised Learning for Estimating $v_\pi(s)$

“training data”:

$$\langle S_0, v_\pi(S_0) \rangle, \langle S_1, v_\pi(S_1) \rangle, \langle S_2, v_\pi(S_2) \rangle \dots$$

We can approximate  $v_\pi(s)$  with

- $G_t$  (MC)
- $R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t)$  (1-step TD)
- $G_{t:(t+n)}$  (n-step TD)
- etc.

-think of these as noisy “measurements” of  $v_\pi(s)$ .

# Supervised Learning for Estimating $q_\pi(s)$

“training data”:

$$\langle (S_0, A_0), q_\pi(S_0, A_0) \rangle, \langle (S_1, A_1), q_\pi(S_1, A_1) \rangle, \langle (S_2, A_2), q_\pi(S_2, A_2) \rangle, \dots$$

We can approximate  $q_\pi(s, a)$  with

- $G_t$  (MC)
- $R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t)$  (1-step SARSA)
- $G_{t:(t+n)}$  (n-step SARSA)
- etc.

-think of these as noisy “measurements” of  $q_\pi(s, a)$ .



# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - **Control Problem via Supervised Learning**
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Supervised Learning for Estimating $q_*(s)$ and $\pi_*(a|s)$

“training data”:

$$\langle (S_0, A_0), q_*(S_0, A_0) \rangle, \langle (S_1, A_1), q_*(S_1, A_1) \rangle, \langle (S_2, A_2), q_*(S_2, A_2) \rangle, \dots$$

We can approximate  $q_*(s, a)$  with

- $G_t$  (MC) and continuously adjust policy
- $R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t)$  (1-step SARSA) and adjust policy
- $G_{t:(t+n)}$  (n-step SARSA) and adjust policy
- $R_{t+1} + \max_a \hat{q}(S_{t+1}, a, \mathbf{w}_t)$  (Q-learning) - no need to adjust policy!
- etc.

-think of these as noisy (and likely very biased in the beginning of the training!) “measurements” of  $q_*(s)$ .

# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Experience Replay

While updating weights  $\mathbf{w}$  using “training data”

$$\langle S_0, v_\pi(S_0) \rangle, \langle S_1, v_\pi(S_1) \rangle, \langle S_2, v_\pi(S_2) \rangle \dots$$

or

$$\langle (S_0, A_0), q_\pi(S_0, A_0) \rangle, \langle (S_1, A_1), q_\pi(S_1, A_1) \rangle, \langle (S_2, A_2), q_\pi(S_2, A_2) \rangle, \dots,$$



the “samples” do not need to be used in the order of time flow. Moreover, one can re-use the “samples”

- this is called *experience replay*.

# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)





# Linear Regression



# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# K Nearest Neighbors (KNN)

KNN with  $k=1$ KNN with  $k=3$ KNN with  $k=5$ KNN with  $k=9$ 




# Contents


- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - **Random Tree / Forest**
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Random Tree / Forest


1 iteration



2 iteration




# Random Tree / Forest



# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - **Logistic Regression**
  - Artificial Neural Networks (NN)

# Logistic Regression



logistic function



$$\sigma(t) = \frac{e^t}{e^t + 1}$$

Decision rule:

if  $x \leq c$  then  $y = 0$  and 1 otherwise

( $c$  is called “decision boundary”)

What if we have 2 inputs:  $x_1$  and  $x_2$ ?







Decision rule:

if point  $(x_1, x_2)$  is “below” the decision boundary then  $y = 0$  and 1 otherwise




# Contents

- 1 Quiz Review
  - Quiz 8
- 2 Approximate Solution Methods (Continued)
  - Mini-batch Gradient Descent
- 3 Reinforcement Learning as Supervised Learning Algorithms
  - Prediction Problem via Supervised Learning
  - Control Problem via Supervised Learning
  - Experience Replay
- 4 Supervised Machine Learning Algorithms: Prediction
  - Linear Regression
  - K Nearest Neighbors (KNN)
  - Random Tree / Forest
  - Logistic Regression
  - Artificial Neural Networks (NN)

# Artificial Neural Networks (NN)



# Artificial Neural Networks (NN)



Convolutional NN (CNN)

