

# CSCI S-89C Deep Reinforcement Learning

Harvard Summer School

Dmitry Kurochkin

Summer 2020

Lecture 11

# Contents

- 1 Quiz Review
  - Quiz 10
- 2 Convolutional Neural Networks (CNN)
  - CNN Layers
    - Convolution
    - Motivation: Visual Cortex
    - Convolutional Layers
    - MaxPooling Layer
    - Example
  - Image Data Augmentation
- 3 Recurrent Neural Networks (RNN)
  - Recurrent Neuron
  - Layer of Recurrent Neurons
  - Memory Cells
  - Long Short-Term Memory (LSTM) Cell
  - Gated Recurrent Unit (GRU) Cell
  - Input and Output Sequences of RNNs

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

- CNN Layers
  - Convolution
  - Motivation: Visual Cortex
  - Convolutional Layers
  - MaxPooling Layer
  - Example

- Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

- Recurrent Neuron
- Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell
- Input and Output Sequences of RNNs

# Quiz 10

## Question 1

4

/ 4 pts

Suppose we want to build a Neural Network for classification. If there are 16 categories to classify, which of the following output layers should be used in Keras?

Please select:

Correct!

- model.add(layers.Dense(16, activation='softmax'))
- model.add(layers.Dense(15, activation='softmax'))
- model.add(layers.Dropout(0.3))
- model.add(layers.Dense(16, activation='relu'))

# Quiz 10

## Question 2

4

/ 4 pts

Suppose we have built the following Neural Network in Keras:

$n = 700$

```
model = models.Sequential()  
model.add(layers.Dense(8, activation='relu', input_shape=(n,)))  
model.add(layers.Dense(8, activation='relu'))  
model.add(layers.Dense(1, activation='relu'))
```

How many parameters does it have?

Correct!

5,689

# Quiz 10

## Question 3

4

/ 4 pts

All of the following activation functions are used in Deep Neural Networks:

$$\sigma^{(1)}(x) = \frac{1}{1+e^{-x}}$$

$$\sigma^{(2)}(x) = \tanh x$$

$$\sigma^{(3)}(x) = \max\{0, x\}$$

$$\sigma^{(4)}(x) = \begin{cases} x, & \text{if } x \geq 0, \\ 0.01x, & \text{otherwise.} \end{cases}$$

Correct!

True

False

# Quiz 10

## Question 4

4

/ 4 pts

Suppose we have the following Neural Network with one input  $x$ , one hidden layer with one neuron  $u$ , and scalar output  $y$ :

$$u = f(w_0^{(1)} + w_1^{(1)}x), \hat{y} = f(w_0^{(2)} + w_1^{(2)}u),$$

where

$$f(x) = \begin{cases} x, & \text{if } x \geq 0, \\ 0, & \text{if } x < 0. \end{cases}$$

The errors are defined as follows:

$$\varepsilon^{(1)} \doteq \frac{\partial L}{\partial u} \text{ and } \varepsilon^{(2)} \doteq \frac{\partial L}{\partial \hat{y}}, \text{ where } L \doteq \frac{1}{2}(\hat{y} - y)^2.$$

If  $(w_0^{(1)}, w_1^{(1)}, w_0^{(2)}, w_1^{(2)})^T = (-1.9, 0.7, 0.3, 1.4)^T$  and we observe  $(x, y) = (1.2, 0.5)$ , please find  $\varepsilon^{(2)}$ .

Correct!

-0.2

# Quiz 10

## Question 5

4

/ 4 pts

Suppose we have the following Neural Network with one input  $x$ , one hidden layer with one neuron  $u$ , and scalar output  $\hat{y}$ :

$$u = f(w_0^{(1)} + w_1^{(1)}x), \hat{y} = f(w_0^{(2)} + w_1^{(2)}u),$$

where

$$f(x) = \begin{cases} x, & \text{if } x \geq 0, \\ 0, & \text{if } x < 0. \end{cases}$$

The errors are defined as follows:

$$\varepsilon^{(1)} \doteq \frac{\partial L}{\partial u} \text{ and } \varepsilon^{(2)} \doteq \frac{\partial L}{\partial \hat{y}}, \text{ where } L \doteq \frac{1}{2}(\hat{y} - y)^2.$$

If  $(w_0^{(1)}, w_1^{(1)}, w_0^{(2)}, w_1^{(2)})^T = (-1.9, 0.7, 0.3, 1.4)^T$  and we observe  $(x, y) = (1.2, 0.5)$ , please find  $\varepsilon^{(1)}$ .

Correct!

-0.28

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

### • CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

- Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

- Recurrent Neuron

- Layer of Recurrent Neurons

- Memory Cells

- Long Short-Term Memory (LSTM) Cell

- Gated Recurrent Unit (GRU) Cell

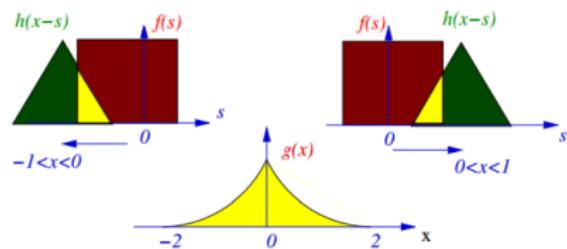
- Input and Output Sequences of RNNs

# Convolution

Convolution of two functions  $f(s)$  and  $h(s)$  is defined as follows:

$$(f * h)(x) \doteq \int_{-\infty}^{+\infty} f(s)h(x - s)ds$$

Example:



# Convolution

Convolution can be used for edge detection. Example:

Original image:

```
plot_image(images[0, :, :, 0])
save_fig("husky-original", tight_layout=False)
plt.show()
```

Saving figure husky-original



# Convolution

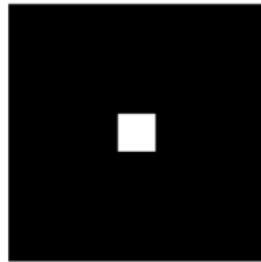
Convolution can be used for edge detection. Example:  
Filter:

```
import numpy as np
import matplotlib.pyplot as plt

def plot_image(image):
    plt.imshow(image, cmap="gray", interpolation="nearest")
    plt.axis("off")

fmap = np.zeros(shape=(7, 7, 1, 3), dtype=np.float32)
fmap[:, 3, 0, 0] = 1
fmap[3, :, 0, 1] = 1
fmap[:, :, 0, 2] = -1
fmap[3, 3, 0, 2] = 48

plot_image(fmap[:, :, 0, 2])
plt.show()
```

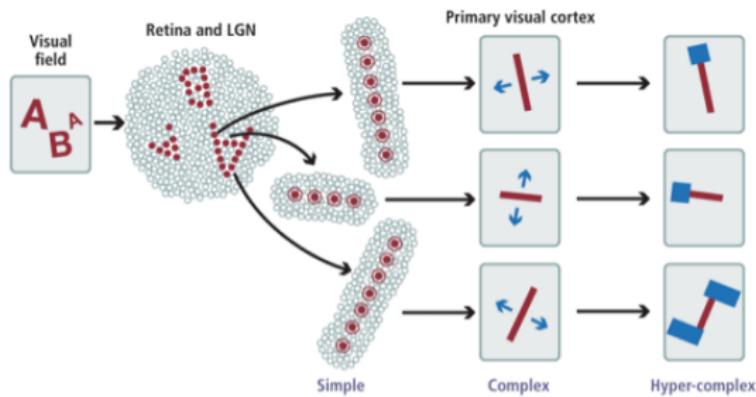


# Convolution

Convolution can be used for edge detection. Example:  
Convolution:

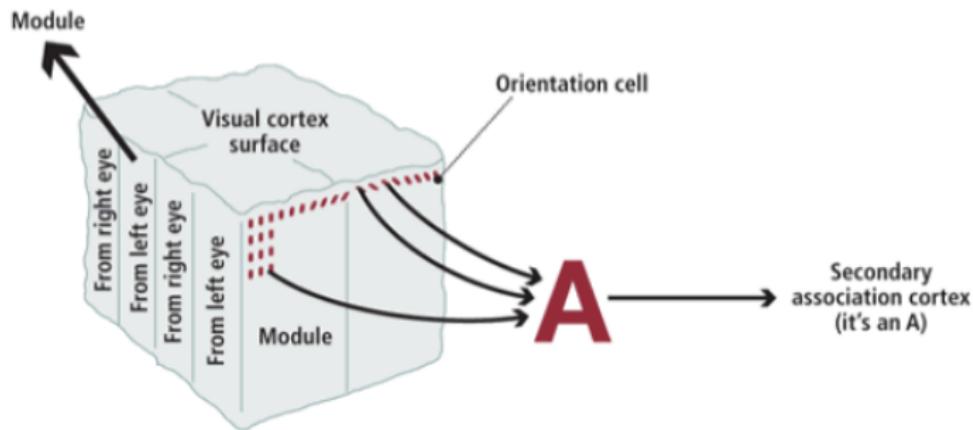


# Visual Cortex

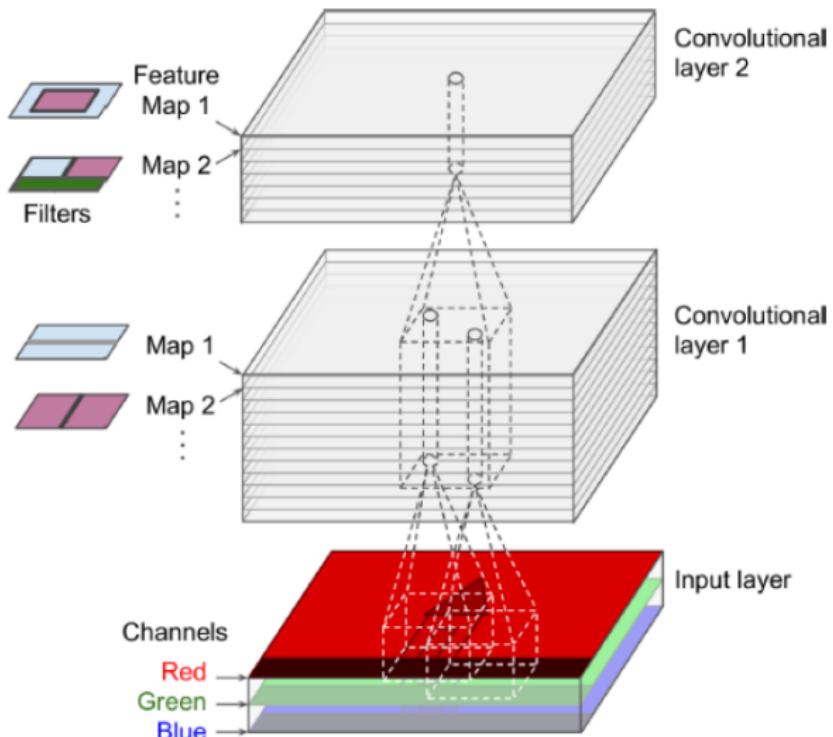


- In the Visual Cortex, simple cells become active when they are subjected to stimuli such as edges.
- Complex cells then combine the information of several simple cells and detect the position and orientation of a structure.
- Hyper complex cells then detect endpoints and crossing lines from this position and orientation information, which is then used in the brain's secondary cortex for information association.

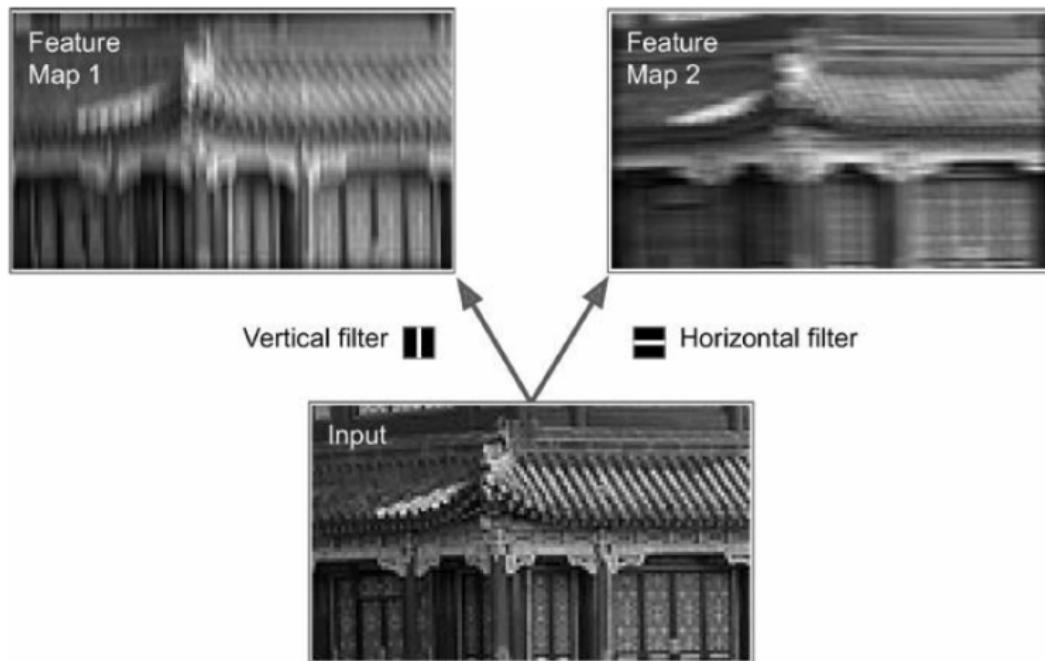
# Visual Cortex



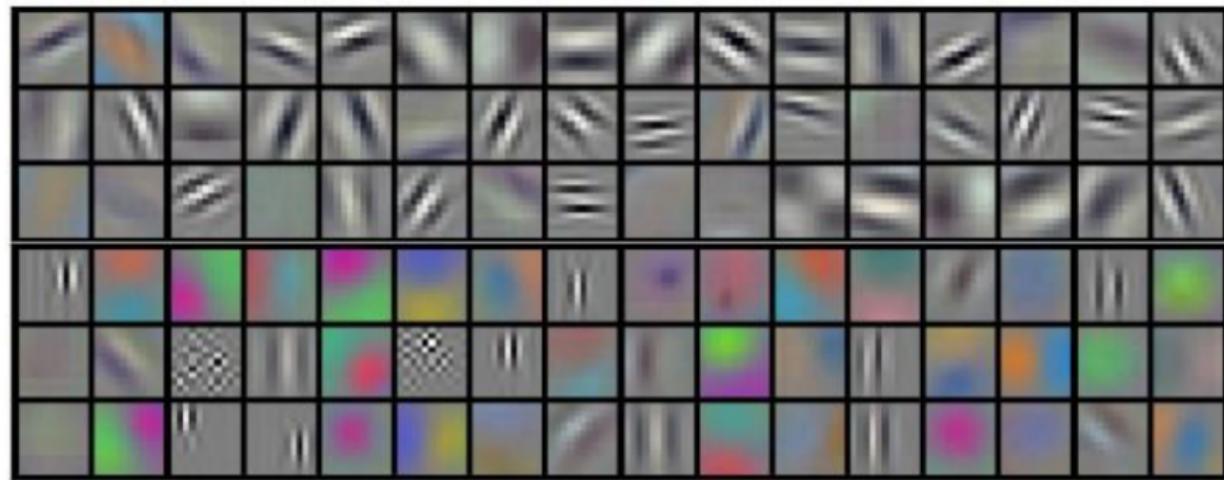
# Convolutional Layers



# Convolutional Layers: Filters

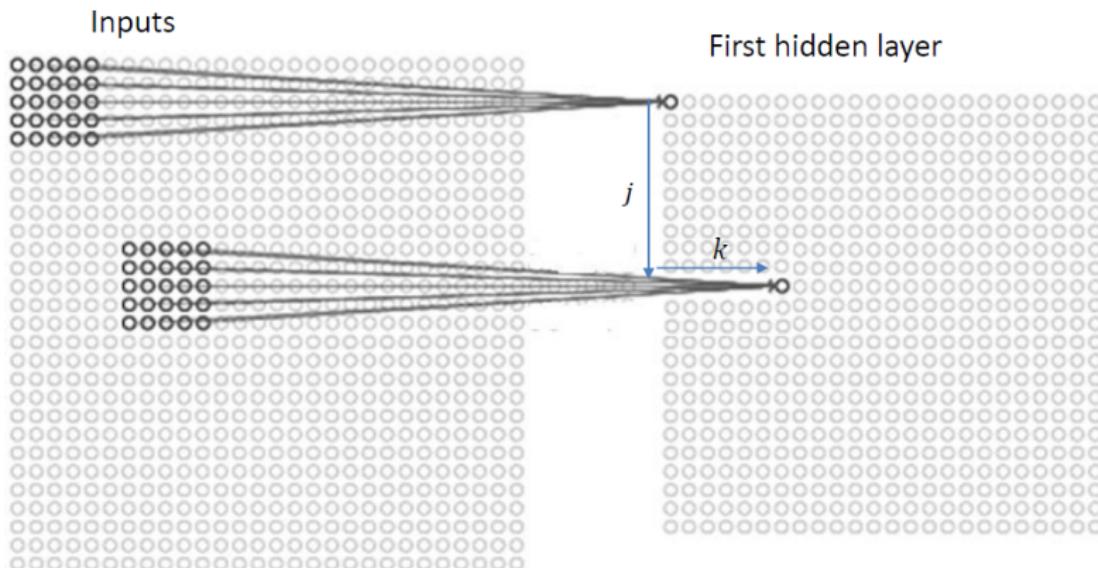


# Convolutional Layers: Filters

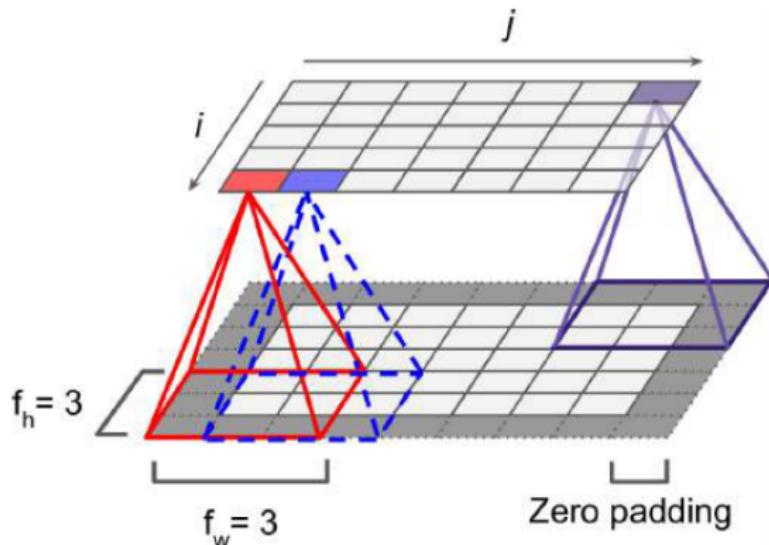


- In early work on CNNs by Krizhevsky et al., they considered 96 filters of size  $11 \times 11$ . Each filter detects certain feature: horizontal edge, vertical edge, slanted edge, etc.
- Eventually, it became apparent that it is not necessary to preselect the filters, but that we could let neural network training process determine the shape of those filters.

# Convolutional Layers: Padding

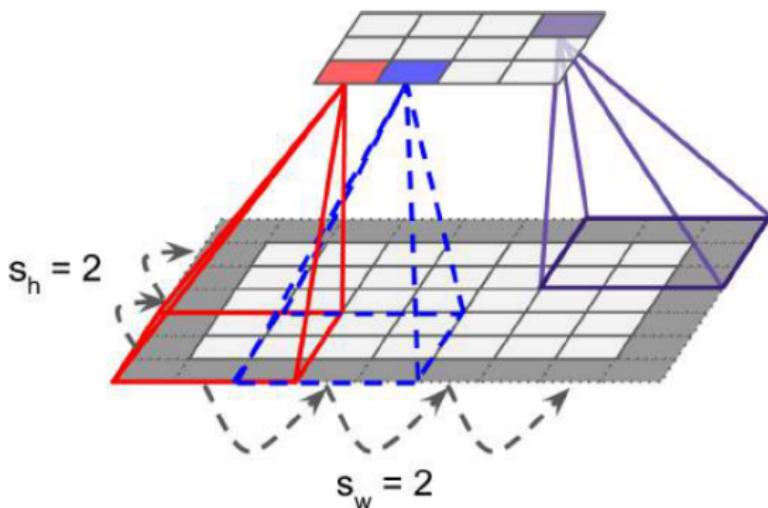


# Convolutional Layers: Padding



Keras: padding can be set to either 'SAME' or 'VALID'

# Convolutional Layers: Strides

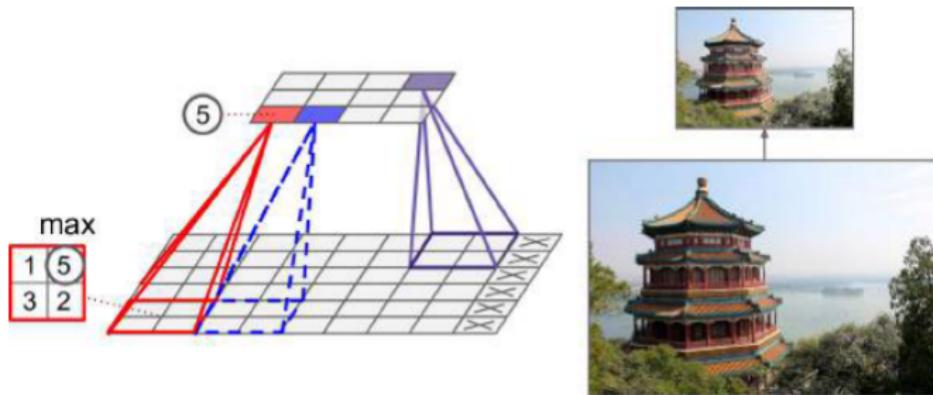


Keras: strides can be set to  $(w_w, w_h)$

# MaxPooling Layer

## Example:

Max pooling layer with  $2 \times 2$  filter, strides= 2, and padding='VALID'



# CNN Example: Keras

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                      input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

- CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

- Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

- Recurrent Neuron

- Layer of Recurrent Neurons

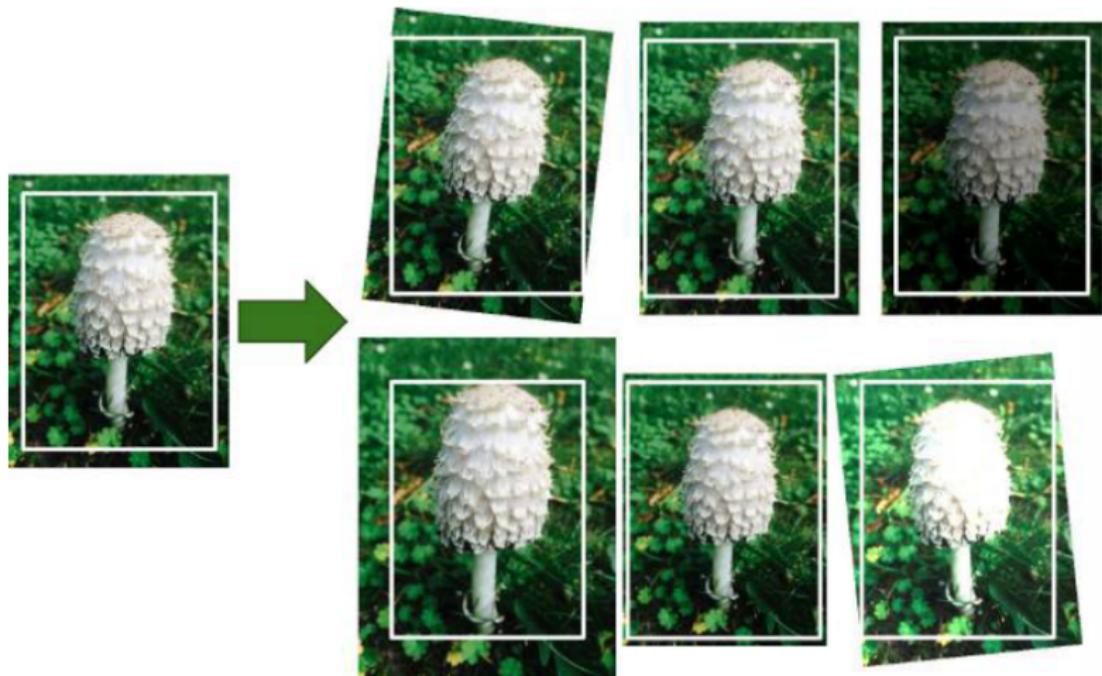
- Memory Cells

- Long Short-Term Memory (LSTM) Cell

- Gated Recurrent Unit (GRU) Cell

- Input and Output Sequences of RNNs

# Image Data Augmentation



# Image Data Augmentation



# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

### CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

### Image Data Augmentation

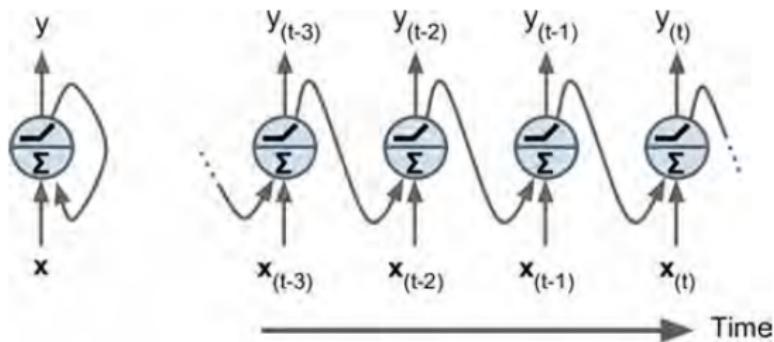
## 3 Recurrent Neural Networks (RNN)

### Recurrent Neuron

- Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell
- Input and Output Sequences of RNNs

# RNN: Recurrent Neuron

Recurrent Neuron:



# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

### CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

### Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

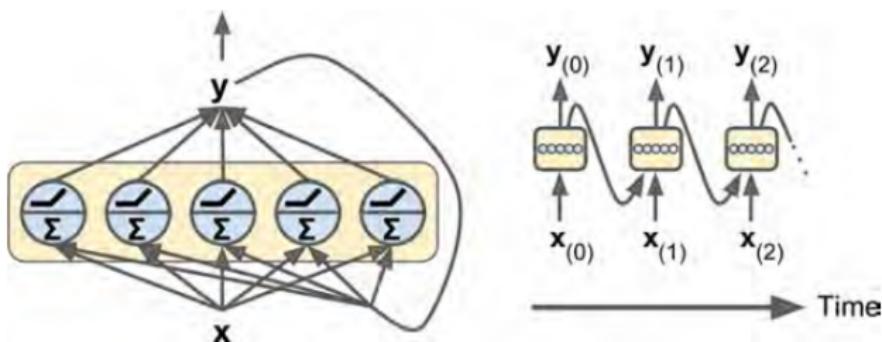
### Recurrent Neuron

### **Layer of Recurrent Neurons**

- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell
- Input and Output Sequences of RNNs

# RNN: Layer of Recurrent Neurons

Layer of Recurrent Neurons:



# Layer of Recurrent Neurons: Keras

Simple RNN in Keras:

```
n_features = 2
n_timesteps = 200

model = models.Sequential()
model.add(layers.SimpleRNN(3, activation='relu', input_shape=(n_timesteps,n_features)))
model.add(layers.Dense(1, activation='linear'))

model.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
simple_rnn_6 (SimpleRNN)	(None, 3)	18
dense_6 (Dense)	(None, 1)	4

Total params: 22

Trainable params: 22

Non-trainable params: 0

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

### CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

### Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

### Recurrent Neuron

### Layer of Recurrent Neurons

### Memory Cells

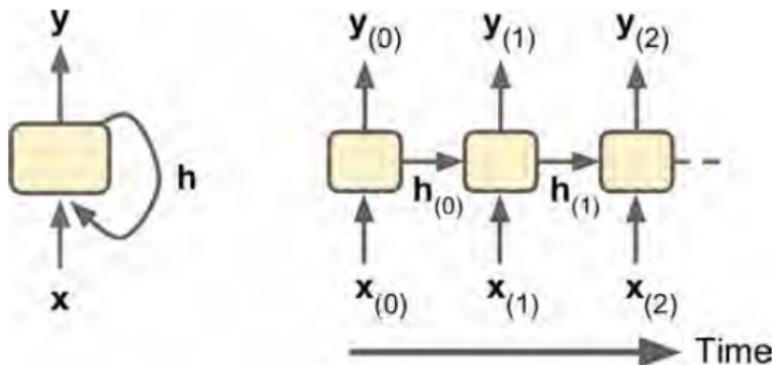
#### Long Short-Term Memory (LSTM) Cell

#### Gated Recurrent Unit (GRU) Cell

#### Input and Output Sequences of RNNs

# RNN: Memory Cells

Layer of Recurrent Neurons:



# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

### CNN Layers

- Convolution
- Motivation: Visual Cortex
- Convolutional Layers
- MaxPooling Layer
- Example

### Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

### Recurrent Neuron

### Layer of Recurrent Neurons

### Memory Cells

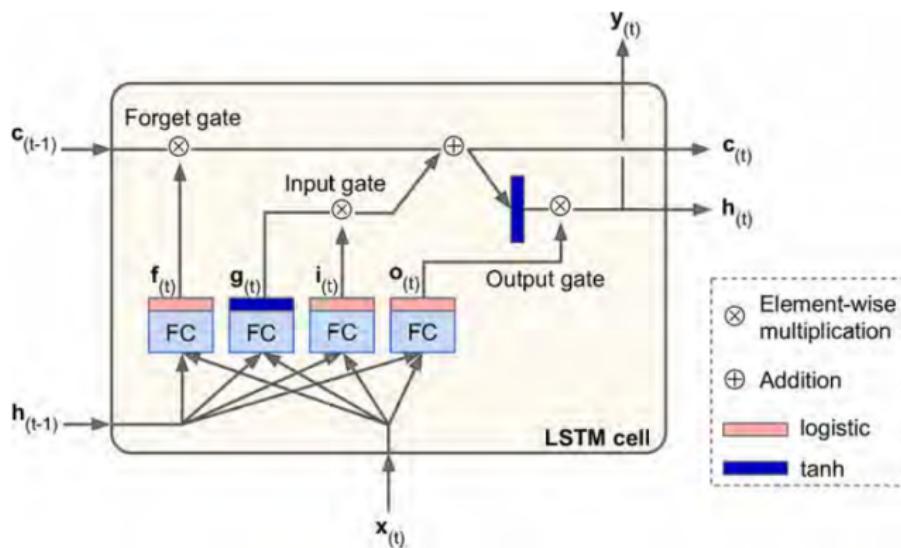
### Long Short-Term Memory (LSTM) Cell

### Gated Recurrent Unit (GRU) Cell

### Input and Output Sequences of RNNs

# Long Short-Term Memory (LSTM) Cell

LSTM Cell:



# Long Short-Term Memory (LSTM) Cell

LSTM Cell Model:

$$\begin{aligned}\mathbf{i}_{(t)} &= \sigma\left(\mathbf{W}_{xi}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hi}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_i\right) \\ \mathbf{f}_{(t)} &= \sigma\left(\mathbf{W}_{xf}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hf}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_f\right) \\ \mathbf{o}_{(t)} &= \sigma\left(\mathbf{W}_{xo}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{ho}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_o\right) \\ \mathbf{g}_{(t)} &= \tanh\left(\mathbf{W}_{xg}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hg}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{b}_g\right) \\ \mathbf{c}_{(t)} &= \mathbf{f}_{(t)} \otimes \mathbf{c}_{(t-1)} + \mathbf{i}_{(t)} \otimes \mathbf{g}_{(t)} \\ \mathbf{y}_{(t)} &= \mathbf{h}_{(t)} = \mathbf{o}_{(t)} \otimes \tanh\left(\mathbf{c}_{(t)}\right)\end{aligned}$$

- $\mathbf{W}_{xi}$ ,  $\mathbf{W}_{xf}$ ,  $\mathbf{W}_{xo}$ ,  $\mathbf{W}_{xg}$  are the weight matrices of each of the four layers for their connection to the input vector  $\mathbf{x}_{(t)}$ .
- $\mathbf{W}_{hi}$ ,  $\mathbf{W}_{hf}$ ,  $\mathbf{W}_{ho}$ , and  $\mathbf{W}_{hg}$  are the weight matrices of each of the four layers for their connection to the previous short-term state  $\mathbf{h}_{(t-1)}$ .
- $\mathbf{b}_i$ ,  $\mathbf{b}_f$ ,  $\mathbf{b}_o$ , and  $\mathbf{b}_g$  are the bias terms for each of the four layers. Note that TensorFlow initializes  $\mathbf{b}_f$  to a vector full of 1s instead of 0s. This prevents forgetting everything at the beginning of training.

# Applications of LSTM Cells

- greatly improved speech recognition on over 4 billion Android phones (since mid 2015)
- greatly improved machine translation through Google Translate (since Nov 2016)
- greatly improved machine translation through Facebook (over 4 billion LSTMbased translations per day as of 2017)
- Siri and Quicktype on almost 2 billion iPhones (since 2016)
- generating answers by Amazon's Alexa and numerous other similar applications.

# Long Short-Term Memory (LSTM) Cell: Keras

LSTM in Keras:

```
n_features = 2
n_timesteps = 200

model = models.Sequential()
model.add(LSTM(16, activation='relu', input_shape=(n_timesteps,n_features)))
model.add(layers.Dense(1, activation='linear'))

model.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 16)	1216

dense_7 (Dense)	(None, 1)	17
-----------------	-----------	----

Total params: 1,233

Trainable params: 1,233

Non-trainable params: 0

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

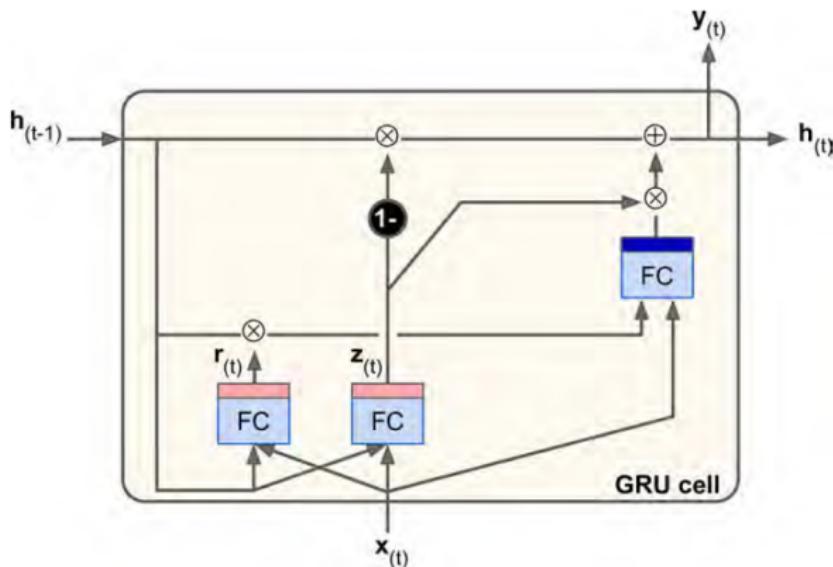
- CNN Layers
  - Convolution
  - Motivation: Visual Cortex
  - Convolutional Layers
  - MaxPooling Layer
  - Example
- Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

- Recurrent Neuron
- Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell
- Input and Output Sequences of RNNs

# Gated Recurrent Unit (GRU) Cell

GRU Cell:



# Gated Recurrent Unit (GRU) Cell

GRU Cell Model:

$$\mathbf{z}_{(t)} = \sigma(\mathbf{W}_{xz}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hz}^T \cdot \mathbf{h}_{(t-1)})$$

$$\mathbf{r}_{(t)} = \sigma(\mathbf{W}_{xr}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hr}^T \cdot \mathbf{h}_{(t-1)})$$

$$\mathbf{g}_{(t)} = \tanh(\mathbf{W}_{xg}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hg}^T \cdot (\mathbf{r}_{(t)} \otimes \mathbf{h}_{(t-1)}))$$

$$\mathbf{h}_{(t)} = (1 - \mathbf{z}_{(t)}) \otimes \tanh(\mathbf{W}_{xg}^T \cdot \mathbf{x}_{(t)} + \mathbf{W}_{hg}^T \cdot \mathbf{h}_{(t-1)} + \mathbf{z}_{(t)} \otimes \mathbf{g}_{(t)})$$

# Contents

## 1 Quiz Review

- Quiz 10

## 2 Convolutional Neural Networks (CNN)

- CNN Layers
  - Convolution
  - Motivation: Visual Cortex
  - Convolutional Layers
  - MaxPooling Layer
  - Example
- Image Data Augmentation

## 3 Recurrent Neural Networks (RNN)

- Recurrent Neuron
- Layer of Recurrent Neurons
- Memory Cells
- Long Short-Term Memory (LSTM) Cell
- Gated Recurrent Unit (GRU) Cell
- Input and Output Sequences of RNNs**

# Input and Output Sequences of RNNs

